

Managerial Brief: MVP Readiness Checklist

What We've Stabilized

- Feature contract: The runtime feature builder now uses a locked schema and scaler so the model always works.
- Pipeline instrumentation: Our stats and sentiment loaders now feed directly into the runtime builder, making it easier to debug.
- Candidate pool workflow: A CLI and HTTP endpoint now produce the deterministic top-30 MVP candidates.
- Documentation & tests: README + agent checklist entries spell out the tooling, while `tests/test\_feature\_builder.py` and `tests/test\_candidate\_pool.py` ensure the tooling works.

What the Plan Requires Next

- Nightly exports: Run `scripts/build\_candidate\_pool\_vectors.py` inside the Python 3.12 sandbox to refresh the candidate pool.
- Endpoint consumption: Point dashboards or clients at `GET /candidate\_pool` with `top\_n`/`pool\_size` query parameters.
- Hand-off ready: Keep the agent checklist and execution plan notes aligned with each run so the next person can pick up.

Risk Brief

- The MLflow/torch artifact demands Python 3.12; running the candidate pool export or the API outside the sandbox may fail.
- If we ever rebuild the artifact under a different interpreter, re-export the scaler/schema files so the runtime can find them.